

A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light green. They are positioned diagonally, with the blue one partially covering the green one.

Practical Linux Kernel Development

Presented by Timothy Finnegan



About Me

- B.S. in Computer Science and Engineering - UNR Spring 2021
- M.S. in Computer Science and Engineering - UNR est. Spring 2023
- One of three primary developers on the university's Nevada Cyber Range cybersecurity education platform.
- Research deals with Android static malware analysis through machine learning.



Linux Kernel Programming

A comprehensive guide to kernel internals, writing kernel modules, and kernel synchronization

Kaiwan N Billimoria





Lesson Outline

- Explore what interacting directly with the Linux kernel looks like
- Describe Linux's "Loadable Kernel Module" (LKM) system.
 - How do you write code that is intended to run within the Kernel, rather than user space?
- What are the Kernel resources available to you from within an LKM?
- A brief outline of some core Kernel Module API functions and functionality.



Major Kernel Subsystems

1. The Core Kernel -- Various topics from this class related to multiprogramming. CPU scheduling, threads, processes, interrupts, etc.
2. Memory Management -- Manages Virtual Address Spaces (VASes) for the Kernel and userspace processes
3. The Virtual Filesystem Switch (VFS) -- The main kernel supports a variety of filesystems (ext4, vfat, ntfs, etc.). The VFS abstracts these filesystems into a single programming interface
4. The Network Protocol Stack -- High-quality implementation of TCP/IP protocols
5. I/O -- Block and Character Devices
6. And more... (Sound, Virtualization, IPC, etc.)

The Kernel as a Monolithic Architecture

- In the Linux Kernel, all Kernel Components live in the same address space (Also known as the Kernel Segment)
- The Kernel Segment cannot be directly accessed by programs in the user space.
- The entire Kernel needs to be compiled as a single program. (time-consuming)
- Any new code compiled with the Kernel will have the same resource access as the Kernel's major subsystems.

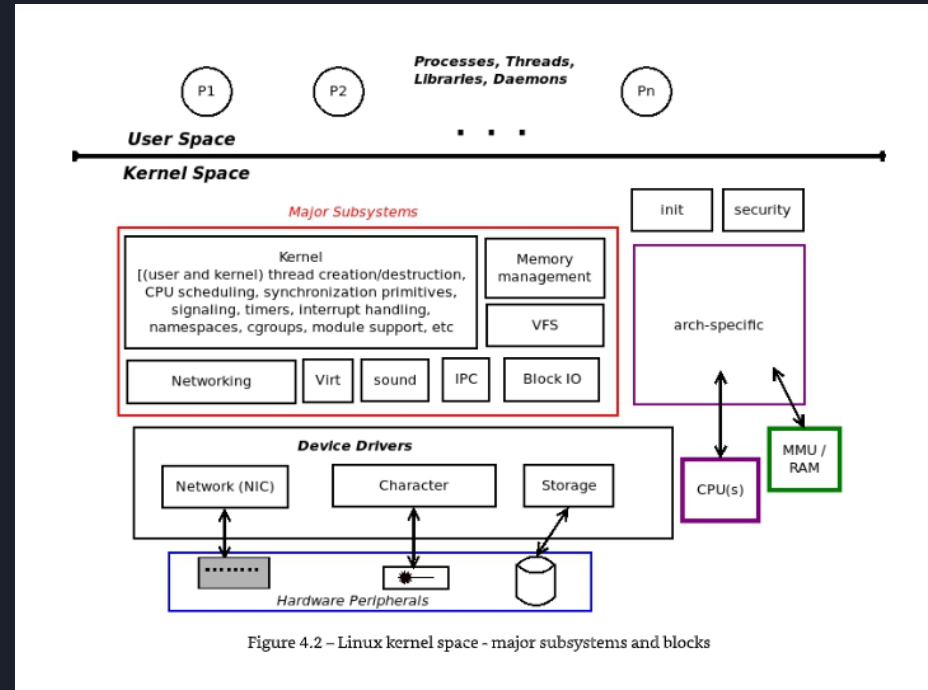


Figure 4.2 – Linux kernel space - major subsystems and blocks



Then how do you add new Kernel functionality (like system drivers) to an existing system?

Loadable Kernel Modules!

Loadable Kernel Modules (LKMs)

- An interface for compiling Kernel code outside of the Kernel's primary source tree.
- Can be loaded into or removed from a system dynamically at runtime.
- Useful for writing hardware-specific code, such as device drivers.
- LKMs exist in the Kernel Space, but do not interact directly with the other major Kernel subsystems.
- Kernel Modules only have access to a subset of the full Kernel API -- Cannot modify some core Kernel functionality (CPU scheduler, timer, or interrupt code).

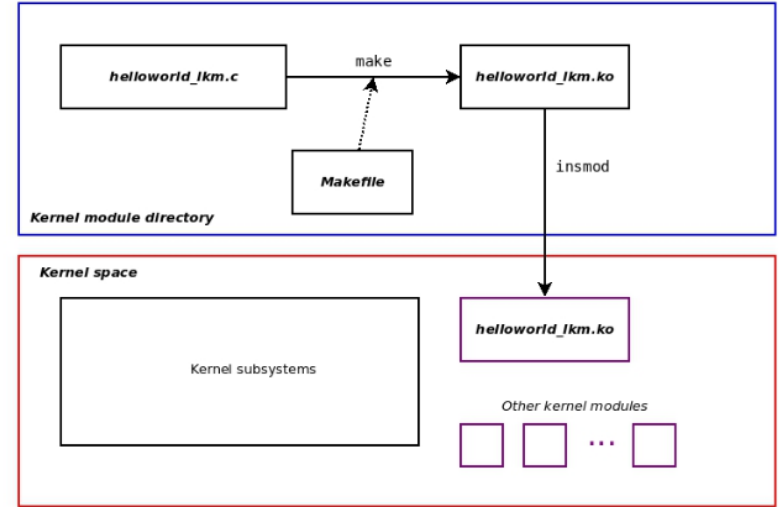


Figure 4.3 – Building and then inserting a kernel module into kernel memory

Worry not: the actual code for both the kernel module C source as well as its Makefile is dealt with in detail in an upcoming section; for now, we want to gain a conceptual understanding only.



How are LKMs loaded into the Kernel?

- LKMs take the form of binary files with the extension *.ko. These files are equivalent to C .o files.
- Kernel modules live in the /lib/modules/[Kernel Version]/kernel/ folder on most Linux systems.
- *insmod* -- “**I**n**s**ert **M**odule”. This is a bash command that takes a Kernel Object File and loads it into the running kernel.
- *rmmod* -- “**R**emove **M**odule”. This is a bash command that takes a running Kernel Module and removes it from the running Kernel.
- With these commands, the Kernel can be modified dynamically at runtime!



Steps to Writing a Kernel Module

1. Download and Compile the desired Linux Kernel version
2. Write your Kernel module.
 - a. Call Module Information Macros
 - b. Define Program Entry Point Function
 - c. Define Program Exit Point (Cleanup) Function
 - d. Set your functions as the default entry and exit points with the appropriate API calls.
3. Compile your Kernel module into a .ko (kernel object) file
4. Place your file into the kernel modules directory
5. Run “insmod” to insert the kernel module into running memory

Kernel Module "Hello, World!"

```
// ch4/helloworld_lkm/helloworld_lkm.c
#include <linux/init.h>
#include <linux/kernel.h>
#include <linux/module.h>
```

```
MODULE_AUTHOR("<insert your name here>");
MODULE_DESCRIPTION("LLKD book:ch4/helloworld_lkm: hello,
world, our first LKM");
MODULE_LICENSE("Dual MIT/GPL");
MODULE_VERSION("0.1");
```

Kernel Macros for
defining module
information.

```
static int __init helloworld_lkm_init(void)
{
    printk(KERN_INFO "Hello, world\n");
    return 0;    /* success */
}
```

The Module Entry Point
(Called by insmod)

Kernel Message Buffer Print Function (Like Printf,
but logs the information to the Kernel)
Accessed with the dmesg command

```
static void __exit helloworld_lkm_exit(void)
{
    printk(KERN_INFO "Goodbye, world\n");
}
```

The Module Exist Point
(Called by rmmod)

Used for Clean-up

```
module_init(helloworld_lkm_init);
module_exit(helloworld_lkm_exit);
```

Macro for setting the entry point function

Macro for setting the exist point function



Common Kernel Module Functions

- **kmalloc(size_t size, gfp_t flags)** -- The Kernel equivalent to malloc. Allocates memory in the Kernel Virtual Address Space
- **kfree(const void * objp)** -- Frees memory allocated by kmalloc
- **cdev_open, cdev_release, cdev_read, cdev_write** -- Character device access functions (Operate on the file representing the character device. Used in device drivers). Similar functions exist for Block devices
- **add_timer** -- Start a timer. Takes a timer_list structure as input
- And more! The LKM API is massive!



Questions?