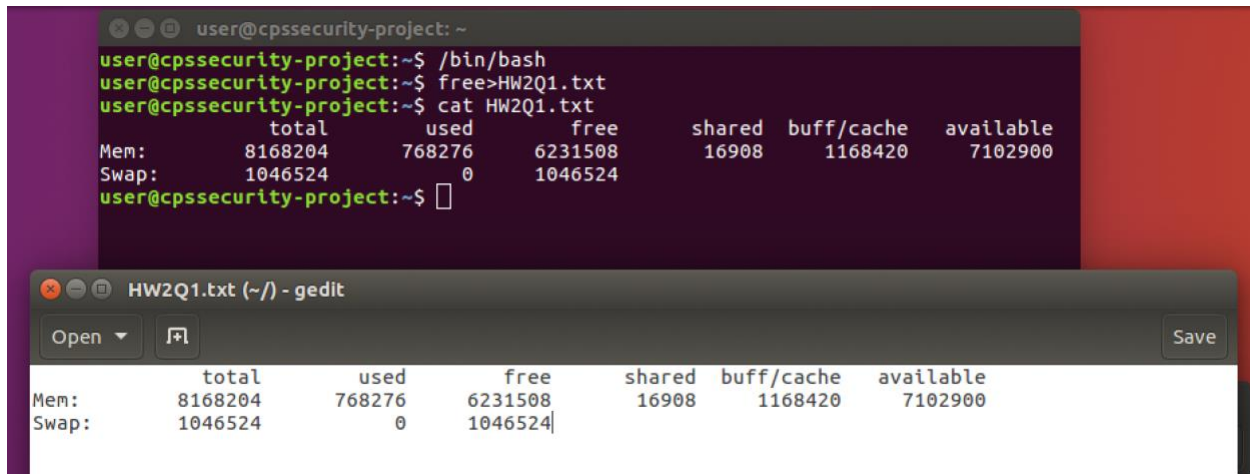


CS 446 Homework 2

Chapter 3

1. [Linux Operating System] free is a command that displays used and available memory in your system. Read man page of free command. Run the command free -o several times, running other programs in between, and store the results in a file. Draw a graph as follows: X-axis: MB-used; for the Y-axis, use (i) Memory Used per unit time; (ii) (Memory Used – Memory Buffered – Memory Cached) per unit time; and (3) Swap Used per unit time. Explain the behavior of this graph with respect to memory utilization in the presence of running various applications.

Answer:



The screenshot shows a terminal window with the following commands and output:


```
user@cpssecurity-project: ~  
user@cpssecurity-project:~$ /bin/bash  
user@cpssecurity-project:~$ free>HW2Q1.txt  
user@cpssecurity-project:~$ cat HW2Q1.txt
```

	total	used	free	shared	buff/cache	available
Mem:	8168204	768276	6231508	16908	1168420	7102900
Swap:	1046524	0	1046524			

The gedit editor window shows the same output as the terminal window.

	Time (sec)	total	used	free	shared	buff/cache	available
Mem:	5	8168204	763984	5883056	16952	1521164	7101804
Mem:	10	8168204	803228	5843220	17352	1521756	7062060
Mem:	15	8168204	818780	5827348	17372	1522076	7046336
Mem:	20	8168204	839048	5807012	17372	1522144	7026036
Mem:	25	8168204	871452	5772184	19112	1524568	6991836
Mem:	30	8168204	1014432	5489884	48892	1663888	6817280
Mem:	35	8168204	1146116	5229216	149604	1792872	6583672
Mem:	40	8168204	1598040	4756548	170188	1813616	6107060
Mem:	45	8168204	1541084	4813244	170184	1813876	6163852
Mem:	50	8168204	1622664	4727468	170248	1818072	6081560
Mem:	55	8168204	1691480	4619892	171404	1856832	6010348
Mem:	60	8168204	1726908	4578532	171936	1862764	5973900

	Time (sec)	total	used	free	shared	buff/cach e	available
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			
Swap:		1046524	0	1046524			



2. [Any System] Plot a histogram and calculate the mean and median of the sizes of the executable binary files on a computer to which you have access. On a Windows system, look at all .exe and .dll files; on a UNIX system look at all executable files in */bin*, */usr/bin*, and */local/bin* that are not scripts (or use the file utility to find all executables). Determine the optimal page size for this computer just considering the code (not data). Consider internal fragmentation and page table size, making some reasonable assumption about the size of a page table entry. Assume that all programs are equally likely to be run and thus should be weighted equally.



Answer:

I have below command to fetch the size of all the executable files (.exe and .dll) in C drive:

EXE: `dir /s /n *.exe | findstr /v .exe.>/Users/mdtamjidh/c_exe.txt`

DLL: `dir /s /n *.dll | findstr /v .dll.>/Users/mdtamjidh/c_dll.txt`

The mean and median of the sizes of the executable binary files are given in the attached *c_exe.xlsx* file (worksheet *exe_2* and *dll_2*)



Chapter 2

3. [Linux Operating System] `df` is a command that displays the amount of disk space available on the file system containing each file name argument. Read man page of `df` command. Run the command `df` to find out how many disk blocks are available and how many are in use. Does the sum of these equals the total number of disk blocks on the disk? If not, explain why there is a difference. Next run the command `df -i` to find out how many i-nodes are available and in use. Now create a new file with just a few characters in it, and again run `df` and `df -i` commands. Explain the effect of creating this new file. Now increase the size of this new file by entering a large number (> 5000) of characters, and again run `df` and `df -i` commands. Explain the effect of increasing the size of the new file.

Answer:

```
user@cpssecurity-project:~$ df
Filesystem      1K-blocks    Used Available Use% Mounted on
udev            4053348         0   4053348   0% /dev
tmpfs           816824     9336   807488   2% /run
/dev/sda1      19478204 15396544   3069180  84% /
tmpfs          4084100        280   4083820   1% /dev/shm
tmpfs           5120          4        5116   1% /run/lock
tmpfs          4084100         0   4084100   0% /sys/fs/cgroup
tmpfs           816824         84   816740   1% /run/user/1000
```

total number of disk blocks = Used + Available

```
user@cpssecurity-project:~$ df -i
Filesystem      Inodes    IUsed   IFree IUse% Mounted on
udev           1013337      415 1012922    1% /dev
tmpfs          1021025      603 1020422    1% /run
/dev/sda1     1245184 532923   712261   43% /
tmpfs          1021025        10 1021015    1% /dev/shm
tmpfs          1021025         5 1021020    1% /run/lock
tmpfs          1021025        17 1021008    1% /sys/fs/cgroup
tmpfs          1021025        39 1020986    1% /run/user/1000
```

total number of disk blocks = Used + Available

After creating a README.txt file and increasing the size of the file (5000M), we can see the changes in the values

```
user@cpssecurity-project:~$ truncate -s 5000M README.txt
user@cpssecurity-project:~$ df
Filesystem      1K-blocks      Used Available Use% Mounted on
udev            4053348         0    4053348   0% /dev
tmpfs           816824       9332    807492   2% /run
/dev/sda1       19478204 15397140    3068584  84% /
tmpfs           4084100        280    4083820   1% /dev/shm
tmpfs           5120           4        5116   1% /run/lock
tmpfs           4084100         0    4084100   0% /sys/fs/cgroup
tmpfs           816824         80    816744   1% /run/user/1000

user@cpssecurity-project:~$ df -i
Filesystem      Inodes    IUsed   IFree IUse% Mounted on
udev            1013337    415   1012922    1% /dev
tmpfs           1021025    601   1020424    1% /run
/dev/sda1       1245184 532936   712248   43% /
tmpfs           1021025     10   1021015    1% /dev/shm
tmpfs           1021025      5   1021020    1% /run/lock
tmpfs           1021025     17   1021008    1% /sys/fs/cgroup
tmpfs           1021025     37   1020988    1% /run/user/1000
```

4. [Programming Problem] Write a program that starts at a given directory and descends the file tree from that point, recording the sizes of all the files it finds. When the traversal is complete, the program should print a histogram of the file sizes using a bin width specified as a parameter into the program (e.g., with 1024, file sizes of 0 to 1023 go in one bin, 1024 to 2047 go in the next, etc).

Answer:

```
user@cpssecurity-project: ~
17678 - 17680 | *
user@cpssecurity-project:~$ ./ceil
Enter the path of directory/directory name: /home/user/Desktop
Enter the size of the bin width: 1024
Histogram of the given directory tree and bin width:
 0 - 1024 | *****
1024 - 2048 | *
2048 - 3072 | *
3072 - 4096 | **
4096 - 5120 | **
5120 - 6144 |
6144 - 7168 |
7168 - 8192 | *
8192 - 9216 |
9216 - 10240 |
10240 - 11264 |
11264 - 12288 |
12288 - 13312 |
13312 - 14336 |
14336 - 15360 |
15360 - 16384 | ***
16384 - 17408 | *
17408 - 18432 |
user@cpssecurity-project:~$
```

Code:

```
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <dirent.h>
#include <errno.h>
#include <fcntl.h>
#include <math.h>
#include <sys/stat.h>
#include <unistd.h>

// define a structure which holds
// stat structure pointer and next pointer
typedef struct fileTreeNode
{
    struct stat *stat_ptr;
    struct fileTreeNode *next;
} StatTreeNode;

// define another structure to hold the list of
// StatTreeNode
typedef struct stat_list
{
    StatTreeNode *start_ptr;
} StatFileList;

// declare a function, that takes the StatFileList initial node
// pointer and returns the maximum file size
off_t getMax_FileSize(StatTreeNode *statTreeNode);

// declare the function that takes a pointer of StatTreeNode and
// string containing the name of the directory and returns the
// int value
int descend_Tree_Dir(StatTreeNode *statTreeNode, const char *pathname);

// declare a function, that takes the StatTreeNode and update the bin array
void upDateBin(StatTreeNode *statTreeNode, int binArray[], int binWidth);

// declare a function to print the histogram
void printHistogram(int binArray[], int num_bins, int bin_width);

// define the main function
int main(int argc, char *argv[])
```

```

{
// declare the stat structure variable
struct stat stat_struct;

// declare the StatFileList variable
StatFileList filesList;

// declare a pointer to the StatTreeNode
StatTreeNode *statTreeNode;
char directoryName[50];

// declare the variable to hold the bin size
int bin_width = 0, num_bins;

// declare a variable to hold the return value of the function
int func_ReturnValue = 0;

// declare a off_t value
off_t max_Size = 0;

// prompt the user for the name of the directory
// directory path name
printf("Enter the path of directory/directory name: ");
scanf("%s", directoryName);

// check the condition whether an object can be created
// with the given directoryName
if (stat(directoryName, &stat_struct) == -1)
{
    perror("The given input is invalid");
    exit(EXIT_FAILURE);
}

// check the condition that whether the given
// directoryName is a directory
if (!S_ISDIR(stat_struct.st_mode))
{
    fputs("The given input path/directory name is not a directory\n", stderr);
    exit(EXIT_FAILURE);
}

// prompt the user for the bin width
printf("Enter the size of the bin width: ");
scanf("%d", &bin_width);

// initialize the first node of the list with

```

```

// initial size of StartTreeNode
filesList.start_ptr = malloc(sizeof(StatTreeNode));

// call the function decend_Tree_Dir() by passing
// initial fileList pointer and name of the directory
func_ReturnValue = decend_Tree_Dir(filesList.start_ptr, directoryName);

// check the condition whether the return value is not zero
// if not zero, exit from the program
if (func_ReturnValue != 0)
{
    exit(EXIT_FAILURE);
}

// set the filesList start_ptr to statTreeNode
statTreeNode = filesList.start_ptr;

// call the function getMax_FileSize() by passing
// the statTreeNode and store the return value in
// max_Size
max_Size = getMax_FileSize(statTreeNode);

// re-set the statTreeNode
statTreeNode = filesList.start_ptr;

// find the number of bins that are required
num_bins = (int)ceil(max_Size / bin_width) + 1;

// declare an array of bins
int bin_array[num_bins];

// initialize each bin to zero value
int i = 0;
for (i = 0; i < num_bins; i++)
{
    bin_array[i] = 0;
}

// call the function uupDateBin, to update the array
// bins
upDateBin(statTreeNode, bin_array, bin_width);

// re-set the statTreeNode
statTreeNode = filesList.start_ptr;

// call the function printHistogram() to diaply the histogram

```



```

    printHistogram(bin_array, num_bins, bin_width);

    return 0;
}

/**
 * printHistogram() function, this accepts an array of int, and two int
 * variables.
 * This function is used to display the histogram of the given bin width.
 */
void printHistogram(int bin_array[], int num_bins, int bin_width)
{
    puts("Histogram of the given directory tree and bin width:");
    for (int i = 0; i < num_bins; i++)
    {
        printf("%5d - %5d\t", bin_width * i, bin_width * (i + 1));
        for (int j = 0; j < bin_array[i]; j++)
        {
            printf("%s", "");
        }
        printf("\n");
    }
}

/**
 * getMax_FileSize() function that accepts a StatTreeNode pointer
 * and returns the off_t.
 * This function, loops through each node in the StatTreeNode(file size),
 * finds the maximum size of the file in the list and returns the maximum
 * size.
 */
off_t getMax_FileSize(StatTreeNode *statTreeNode)
{
    // declare a variable of type off_t
    off_t largeFileSize = 0;

    // define a variable pointer of StatTreeNode, which is
    // initialized with the statTreeNode
    StatTreeNode *current = statTreeNode;

    // loop through each node
    while (current->next != NULL && current->stat_ptr != NULL)
    {
        // condition to check the largest file size
        if (current->stat_ptr->st_size > largeFileSize)
        {

```

```

        // if the current's stat size is larger than the
        // largeFileSize then set the current's stat size
        // to largeFileSize
        largeFileSize = current->stat_ptr->st_size;
    }

    // move pointer to the next node
    current = current->next;
}

// return the largeFileSize
return largeFileSize;
}

/**
 * upDateBin() function that accepts a StatTreeNode pointer, an int array
 * and an int value
 */
void upDateBin(StatTreeNode *statTreeNode, int binArray[], int binWidth)
{
    // define a variable pointer of StatTreeNode, which is
    // initialized with the statTreeNode
    StatTreeNode *current = statTreeNode;

    // declare a variable to hold the index
    off_t index = 0;
    // loop through each node
    while (current->next != NULL && current->stat_ptr != NULL)
    {
        // set the index value
        index = current->stat_ptr->st_size / binWidth;

        // increment the binArray value at the index
        binArray[index]++;

        // move pointer to the next node
        current = current->next;
    }
}

/**
 * decend_Tree_Dir() recursive function accepts a StatTreeNode and a string
 * holding path name and returns an int value.
 * This function goes through each directory and file and sets the
 * size of the file to the statTreeNode node and returns an int value
 * if there is any error.

```

```

*/
int descend_Tree_Dir(StatTreeNode *statTreeNode, const char *directory_PathName)
{
    // declare the variable of type DIR
    DIR *directoryInput;

    // declare an int variable
    int dir_fd;

    // declare the variable to hold the file status
    int file_status = -1;

    // declare a variable to hold the error status
    int err_status = -1;

    // declare the dirent structure pointer
    struct dirent *direntPtr;

    // declare the stat structure pointer
    struct stat *stat_buffer;

    // Check whether the given directory name is able to open or not
    if ((directoryInput = opendir(directory_PathName)) == NULL)
    {
        fprintf(stderr, "Unable to open \"%s\" directory.\n", directory_PathName);
        return errno;
    }

    // Check whether the given directory is able to open the file descriptor
    if ((dir_fd = dirfd(directoryInput)) == -1)
    {
        fprintf(stderr, "Could not able to obtain directory file descriptor "
            "of: %s\n", directory_PathName);
        return errno;
    }

    // if able to open up the directory then loop through and get the size of the file
    while ((direntPtr = readdir(directoryInput)) != NULL)
    {
        // if the directory name contains "." or ".." then skip the directories
        if (strcmp(direntPtr->d_name, ".") == 0 || strcmp(direntPtr->d_name, "..") == 0)
        {
            continue;
        }

        // initialize the stat_buffer

```

```

stat_buffer = malloc(sizeof(struct stat));

// get the status of the file
file_status = fstatat(dir_fd, direntPtr->d_name, stat_buffer, 0);

// if the file_status is -1, then display an error message and return
// error value
if (file_status == -1)
{
    // get the error number
    err_status = errno;

    // free the stat_buffer
    free(stat_buffer);

    // display the error message
    fprintf(stderr, "Unable to get the file status related to "
        "the file \"%s\" descriptor. \n", direntPtr->d_name);

    // return the error status
    return err_status;
}

// use switch structure to invoke the related case
switch (stat_buffer->st_mode & S_IFMT)
{
    // if the file is a regular file, then set the size of the
    // to the
case S_IFREG:
    statTreeNode->next = malloc(sizeof(StatTreeNode));
    statTreeNode->stat_ptr = stat_buffer;
    statTreeNode = statTreeNode->next;
    statTreeNode->next = NULL;
    continue;

    // if the file is a directory with in a directory, the
    // go through each sub-directory for the files by calling
    // current function(recursive)
case S_IFDIR:
{
    // add the path
    char *sub_path = malloc(strlen(directory_PathName) + strlen(direntPtr->d_name)
+ 2);

    // display the subpath
    sprintf(sub_path, "%s/%s", directory_PathName, direntPtr->d_name);

```

```

// get the error status returned by the function decend_Tree_Dir()
err_status = decend_Tree_Dir(statTreeNode, sub_path);

//
if (err_status != 0)
{
    fprintf(stderr, "Unable to open the sub_directory: %s\n", direntPtr->d_name);
    return err_status;
}

// free the pointer
free(sub_path);
break;
}
}

// free the stat_buffer
free(stat_buffer);
}

// close the directory
if (closedir(directoryInput) == -1)
{
    fprintf(stderr, "Could not close the \"%s\" directory.\n", directory_PathName);
    return errno;
}
return 0;
}

```